

What is MapReduce

- ❖ Origin from Google, [OSDI'04]
 - MapReduce: Simplified Data Processing on Large Clusters
 - Jeffrey Dean and Sanjay Ghemawat
- ❖ Programming model for parallel data processing
- ❖ Hadoop can run MapReduce programs written in various languages: e.g. Java, Ruby, Python, C++
- ❖ For large-scale data processing
 - Exploits large set of commodity computers
 - Executes process in a distributed manner
 - Offers high availability

Motivation for MapReduce

- ❖ There was need for an abstraction that hides many system-level details from the programmer.
- ❖ MapReduce addresses this challenge by providing a simple abstraction for the developer, transparently handling most of the details behind the scenes in a **scalable**, **robust**, and **efficient** manner.
- ❖ MapReduce separates the **what** from the **how**

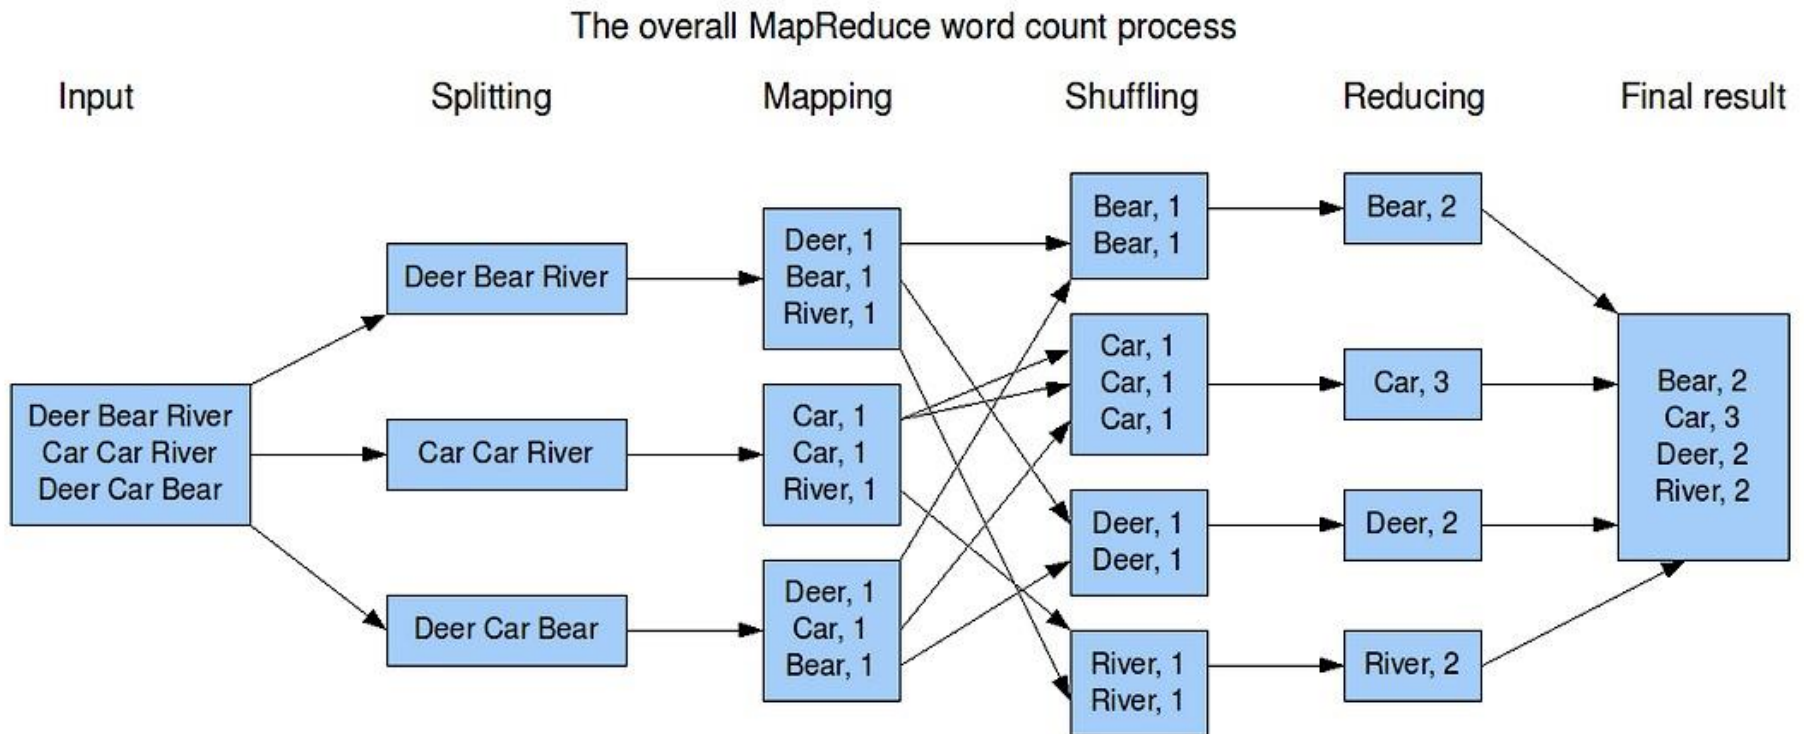
The Idea of MapReduce

- ❖ Inspired by the map and reduce functions in functional programming
- ❖ We can view map as a transformation over a dataset
 - This transformation is specified by the function f
 - Each functional application happens in **isolation**
 - The application of f to each element of a dataset can be parallelized in a straightforward manner
- ❖ We can view reduce as an aggregation operation
 - The aggregation is defined by the function g
 - Data locality: elements in the list must be “brought together”
 - If we can **group** elements of the list, also the reduce phase can proceed in parallel
- ❖ The framework coordinates the map and reduce phases:
 - Grouping intermediate results happens in parallel

Everything Else?

- ❖ Handles scheduling
 - Assigns workers to map and reduce tasks
- ❖ Handles “data distribution”
 - Moves processes to data
- ❖ Handles synchronization
 - Gathers, sorts, and shuffles intermediate data
- ❖ Handles errors and faults
 - Detects worker failures and restarts
- ❖ Everything happens on top of a distributed file system (HDFS)
- ❖ You don't know:
 - Where mappers and reducers run
 - When a mapper or reducer begins or finishes
 - Which input a particular mapper is processing
 - Which intermediate key a particular reducer is processing

MapReduce Example - WordCount



- ❖ Hadoop MapReduce is an implementation of MapReduce
 - MapReduce is a computing paradigm (Google)
 - Hadoop MapReduce is an open-source software

Data Structures in MapReduce

- ❖ Key-value pairs are the basic data structure in MapReduce
 - Keys and values can be: integers, float, strings, raw bytes
 - They can also be arbitrary data structures, but must be comparable (for sorting)
- ❖ The design of MapReduce algorithms involves:
 - Imposing the key-value structure on arbitrary datasets
 - ▶ E.g.: for a collection of Web pages, input keys may be URLs and values may be the HTML content
 - In some algorithms, input keys are not used (e.g., wordcount), in others they uniquely identify a record
 - Keys can be combined in complex ways to design various algorithms

Map and Reduce Functions


- ❖ Programmers specify two functions:
 - **map** $(k_1, v_1) \rightarrow \text{list } [<k_2, v_2>]$
 - ▶ Map transforms the input into key-value pairs to process
 - **reduce** $(k_2, \text{list } [v_2]) \rightarrow [<k_3, v_3>]$
 - ▶ Reduce aggregates the list of values for each key
 - ▶ All values with the same key are sent to the same reducer
 - $\text{list } [<k_2, v_2>]$ will be grouped according to key k_2 as $(k_2, \text{list } [v_2])$
- ❖ The MapReduce environment takes in charge of everything else...
- ❖ A complex program can be decomposed as a succession of Map and Reduce tasks

Understanding MapReduce

❖ Map>>

- $(K1, V1) \rightarrow$
 - ▶ Info in
 - ▶ Input Split
- $\text{list}(K2, V2)$
 - ▶ Key / Value out (intermediate values)
 - ▶ One list per local node
 - ▶ Can implement local Reducer (or Combiner)

Shuffle/Sort>>



Reduce

- $(K2, \text{list}(V2)) \rightarrow$
 - ▶ Shuffle / Sort phase precedes Reduce phase
 - ▶ Combines Map output into a list
- $\text{list}(K3, V3)$
 - ▶ Usually aggregates intermediate values

$(input) <k1, v1> \rightarrow \text{map} \rightarrow <k2, v2> \rightarrow \text{combine} \rightarrow <k2, \text{list}(V2)> \rightarrow \text{reduce} \rightarrow <k3, v3> (output)$

WordCount – Mapper/Reducer

❖ Let's count number of each word in documents (e.g., Tweets/Blogs)

➤ Reads input pair $\langle k1, v1 \rangle$

▶ The input to the mapper is in format of $\langle docID, docText \rangle$:

$\langle D1, \text{"Hello World"} \rangle, \langle D2, \text{"Hello Hadoop Bye Hadoop"} \rangle$

➤ Outputs pairs $\langle k2, v2 \rangle$

▶ The output of the mapper is in format of $\langle term, 1 \rangle$:

$\langle \text{Hello}, 1 \rangle \langle \text{World}, 1 \rangle \langle \text{Hello}, 1 \rangle \langle \text{Hadoop}, 1 \rangle \langle \text{Bye}, 1 \rangle \langle \text{Hadoop}, 1 \rangle$


➤ After shuffling and sort, reducer receives $\langle k2, \text{list}(v2) \rangle$

$\langle \text{Hello}, \{1, 1\} \rangle \langle \text{World}, \{1\} \rangle \langle \text{Hadoop}, \{1, 1\} \rangle \langle \text{Bye}, \{1\} \rangle$

➤ The output is in format of $\langle k3, v3 \rangle$:

$\langle \text{Hello}, 2 \rangle \langle \text{World}, 1 \rangle \langle \text{Hadoop}, 2 \rangle \langle \text{Bye}, 1 \rangle$

A Brief View of MapReduce



Shuffle and Sort


❖ Shuffle

- Input to the Reducer is the sorted output of the mappers. In this phase the framework fetches the relevant partition of the output of all the mappers, via HTTP.

❖ Sort

- The framework groups Reducer inputs by keys (since different Mappers may have output the same key) in this stage.

❖ Hadoop framework handles the Shuffle and Sort step .



“Hello World” in MapReduce

❖ Input:

- Key-value pairs: (docid, doc) of a file stored on the distributed filesystem
- docid : unique identifier of a document
- doc: is the text of the document itself

❖ Mapper:

- Takes an input key-value pair, tokenize the line
- Emits intermediate key-value pairs: the word is the key, and the integer is the value

❖ The framework:

- Guarantees all values associated with the same key (the word) are brought to the same reducer

❖ The reducer:

- Receives all values associated to some keys
- Sums the values and writes output key-value pairs: the key is the word, and the value is the number of occurrences